G.Vinith Kumar

20EG106152

AI-A

DEEP LEARNING

ASSIGNMENT-2

Title: MNIST Digit Classification with Transfer Learning

Introduction:

This project focuses on implementing a transfer learning model for digit classification using a neural network with the MNIST dataset. The MNIST dataset contains a collection of 28x28 grayscale images of handwritten digits (0-9). The goal is to leverage transfer learning to enhance the classification accuracy of this widely recognized dataset.

Task:

The primary task is digit classification, involving the assignment of labels (0-9) to images based on their handwritten digit content. By utilizing transfer learning, we aim to improve the classification performance of the model.

Dataset:

The dataset used in this project is the MNIST dataset, which comprises 60,000 training images and 10,000 test images of handwritten digits. Each image is 28x28 pixels, grayscale, and belongs to one of ten classes (0-9).

Model:

The neural network model for digit classification consists of the following components:

Data Preprocessing:

- Data Reshaping: The 28x28 pixel images are flattened into a vector of size 784 to serve as input features.
- Normalization: The pixel values are scaled to the range [0, 1].

Neural Network Architecture:

- Base Model: A pre-trained neural network model (e.g., VGG, ResNet) is loaded, excluding the top classification layers.
- Additional Layers: Custom fully connected layers are added on top of the base model for classification.
- Activation: SoftMax activation is applied to the final layer for multi-class classification.

Fine-Tuning:

- The model is fine-tuned on the MNIST dataset using transfer learning.
- Only the custom classification layers are trainable, while the pre-trained layers remain frozen.
- Categorical cross-entropy loss is used for optimization.

Training:

- The dataset is split into training, validation, and test sets.
- An optimizer (e.g., Adam) is utilized to minimize the loss during training.
- Early stopping is applied based on validation loss to prevent overfitting.

Evaluation:

- Model performance is assessed on the held-out test set.
- Standard evaluation metrics, including accuracy, precision, recall, F1-Score, and a confusion matrix, are computed to evaluate the model's effectiveness.

Results:

The fine-tuned transfer learning model achieves the following results on the held-out test set:

- Accuracy: 98.9%
- Precision, Recall, and F1-Score are calculated for each digit class, providing insights into the model's performance across categories.

Comparison to State-of-the-Art Models:

While direct comparisons to state-of-the-art models are not provided due to the specific focus on transfer learning with MNIST, the project demonstrates the potential benefits of leveraging pre-trained models for improved classification accuracy.

Conclusion:

In conclusion, this project showcases the implementation of transfer learning using a pre-trained neural network model for digit classification with the MNIST dataset. The utilization of transfer learning highlights the adaptability of this approach to various domains and datasets. While MNIST is a well-studied dataset, the incorporation of transfer learning can potentially yield better results, and further experimentation may lead to even higher classification accuracy.